

The background of the slide features several overlapping, curved, metallic-looking shapes that resemble stylized 'C' or 'G' characters. These shapes are rendered with a brushed metal texture and are set against a dark, fine-grained background. The lighting creates highlights and shadows, giving the shapes a three-dimensional appearance.

Accelerating Linpack with CUDA on heterogeneous clusters

Massimiliano Fatica
mfatica@nvidia.com



Outline



- Linpack benchmark
- CUBLAS and DGEMM
- Results:
 - Accelerated Linpack on workstation
 - Accelerated Linpack on heterogeneous cluster

Linpack



LINPACK Benchmark

The LINPACK benchmark is very popular in the HPC space, because it is used as a performance measure for ranking supercomputers in the TOP500 list.

The most widely used implementation is the HPL software package from the Innovative Computing Laboratory at the University of Tennessee:

it solves a random dense linear system in double precision arithmetic on distributed-memory computers.

LINPACK Benchmark



Solve a dense NxN linear system:


$$Ax=b$$

Solution is obtained by Gaussian elimination with partial pivoting

Floating point workload:

$$\frac{2}{3} N^3 + 2 N^2$$

(LU decomposition) (back solve)



Factorize the current block (red), update the green and yellow parts when done

The bigger the problem size N is, the more time is spent in the update of the trailing matrices (DGEMM)

CUDA Accelerated LINPACK



Both CPU cores and GPUs are used in synergy with minor or no modifications to the original source code (HPL 2.0):


- An host library intercepts the calls to DGEMM and DTRSM and executes them simultaneously on the GPUs and CPU cores. Library is implemented with CUBLAS
- Use of pinned memory for fast PCI-e transfers, up to 5.7GB/s on x16 gen2 slots. Only changes to the HPL source

CUBLAS and DGEMM

CUBLAS

- **Implementation of BLAS (Basic Linear Algebra Subprograms) on top of CUDA driver**
 - Self-contained at the API level, no direct interaction with CUDA driver
- **Basic model for use**
 - Create matrix and vector objects in GPU memory space
 - Fill objects with data
 - Call sequence of CUBLAS functions
 - Retrieve data from GPU
- **CUBLAS library contains helper functions**
 - Creating and destroying objects in GPU space
 - Writing data to and retrieving data from objects
 - Error handling
- **Function naming convention**
 - cublas + BLAS name: eg., cublasDGEMM
 - Easier to mix CPU and GPU BLAS calls

DGEMM: $C = \alpha A B + \beta C$



$$\text{DGEMM}(A, B, C) = \text{DGEMM}(A, B_1, C_1) \cup \text{DGEMM}(A, B_2, C_2)$$

(GPU) (CPU)

The idea can be extended to multi-GPU configuration and to handle huge matrices

Find the optimal split, knowing the relative performances of the GPU and CPU cores on DGEMM



Overlap DGEMM on CPU and GPU

Copy A from CPU memory to GPU memory devA

```
status = cublasSetMatrix (m, k, sizeof(A[0]), A, lda, devA, m_gpu);
```

Copy B1 from CPU memory to GPU memory devB

```
status = cublasSetMatrix (k, n_gpu, sizeof(B[0]), B, ldb, devB, k_gpu);
```

Copy C1 from CPU memory to GPU memory devC

```
status = cublasSetMatrix (m, n_gpu, sizeof(C[0]), C, ldc, devC, m_gpu);
```

Perform DGEMM(devA,devB,devC) on GPU

Control immediately return to CPU

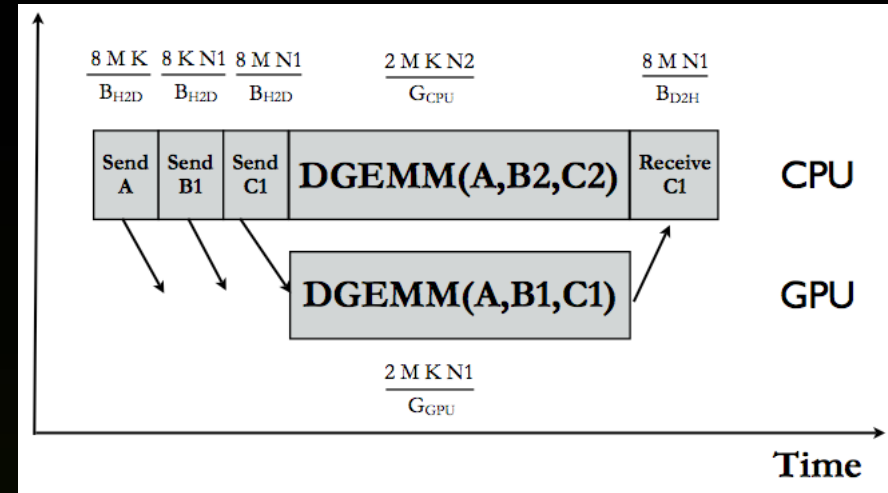
```
cublasDgemm('n', 'n', m, n_gpu, k, alpha, devA, m, devB, k, beta, devC, m);
```

Perform DGEMM(A,B2,C2) on CPU

```
ldgemm('n','n',m,n_cpu,k, alpha, A, lda,B+ldb*n_gpu, ldb, beta,C+ldc*n_gpu, ldc);
```

Copy devC from GPU memory to CPU memory C1

```
status = cublasGetMatrix (m, n, sizeof(C[0]), devC, m, C, *ldc);
```



Using CUBLAS, it is very easy to express the workflow in the diagram

DGEMM performance on GPU



A DGEMM call in CUBLAS maps to several different kernels depending on the size of the matrices. With the combined CPU/GPU approach, we can always send optimal work to the GPU.

M	K	N	M%64	K%16	N%16	Gflops
448	400	12320	Y	Y	Y	82.4
12320	400	1600	N	Y	Y	75.2
12320	300	448	N	N	Y	55.9
12320	300	300	N	N	N	55.9

Tesla T10 1.44Ghz, data resident in GPU memory. Optimal kernel achieves 95% of peak

Optimal split




If $A(M,K)$, $B(K,N)$ and $C(M,N)$, a DGEMM call performs $2 \cdot M \cdot K \cdot N$ operations


$$T_{\text{CPU}}(M,K,N_2) = T_{\text{GPU}}(M,k,N_1) \quad N=N_1+N_2$$

If G_{CPU} denotes the DGEMM performance of the CPU in Gflops and G_{GPU} the one of the GPU,
The optimal split is

$$\eta = G_{\text{GPU}} / (G_{\text{CPU}} + G_{\text{GPU}})$$



DGEMM Performance



Results



Results on workstation

SUN Ultra 24 workstation with an Intel Core2 Extreme Q6850 (3.0Ghz) CPU, 8GB of memory plus a Tesla C1060 (1.296Ghz) card.

- Peak DP CPU performance of 48 GFlops
- Peak DP GPU performance of 77 GFlops (60*clock)

T/V	N	NB	P	Q	Time	Gflops
WR00L2L2	23040	960	1	1	97.91	8.328e+01
$\ Ax-b\ _{\infty}/(\epsilon * (\ A\ _{\infty} * \ x\ _{\infty} + \ b\ _{\infty}) * N) =$						0.0048141 PASSED

83.2 Gflops sustained on a problem size using less than 4GB of memory: 66% of efficiency

T/V	N	NB	P	Q	Time	Gflops
WR00C2R2	32320	1152	1	1	246.86	9.118e+01
$\ Ax-b\ _{\infty}/(\epsilon * (\ A\ _{\infty} * \ x\ _{\infty} + \ b\ _{\infty}) * N) =$						0.0042150 PASSED

91.1 Gflops sustained on a problem size using 8GB of memory: 73% of efficiency

PCI-e transfer speed

	SUN Ultra 24 PCI-e x16 gen2		Supermicro 6015TW PCI-e x16 gen2	
	Pageable Memory	Pinned Memory	Pageable Memory	Pinned Memory
Host to Device	2132 MB/s	5212 MB/s	2524 MB/s	5651 MB/s
Device to Host	1882 MB/s	5471 MB/s	2084 MB/s	5301 MB/s

CUDA offers a fast PCI-e transfer when host memory is allocated with `cudaMallocHost` instead of regular `malloc`. Limited to 4GB in CUDA 2.1, no limitations in upcoming CUDA 2.2



T/V	N	NB	P	Q	Time	Gflops
WR00L2L2	23040	960	1	1	97.91	8.328e+01
Ax-b _oo/(eps*(A _oo* x _oo+ b _oo)*N)=						0.0048141 PASSED

T/V	N	NB	P	Q	Time	Gflops
WR00L2L2	23040	960	1	1	117.06	69.66e+01
Ax-b _oo/(eps*(A _oo* x _oo+ b _oo)*N)=					0.0048141 PASSED	



Results on cluster

Cluster with 8 nodes, each node connected to half of a Tesla S1070-500 system:

- Each node has 2 Intel Xeon E5462 (2.8Ghz with 1600Mhz FSB) , 16GB of memory and 2 GPUs (1.44Ghz clock).
- The nodes are connected with SDR Infiniband.

T/V	N	NB	P	Q	Time	Gflops
WR10L2R4	92164	960	4	4	479.35	1.089e+03
$\ Ax-b\ _{\infty}/(\epsilon \cdot (\ A\ _{\infty} \cdot \ x\ _{\infty} + \ b\ _{\infty}) \cdot N) =$					0.0026439 PASSED

1.089 Tflop/s sustained using less than 4GB of memory per MPI process.

The first system to break the Teraflop barrier was ASCI Red (1.068 Tflop/s) in June 1997 with 7264 Pentium Pro processors. The GPU accelerated nodes occupies 8U (11U including Ethernet and Infiniband switches).

Results on cluster

Increasing the problem size (8GB per MPI process, CUDA 2.2 beta):

T/V	N	NB	P	Q	Time	Gflops
WR11R2L2	118144	960	4	4	874.26	1.258e+03
$\ Ax-b\ _{\infty}/(\text{eps}*(\ A\ _{\infty}*\ x\ _{\infty}+\ b\ _{\infty})*N)=$					0.0031157	PASSED

Results on cluster



11/30/2008

71

Computer (Full Precision)	Number of Procs or Cores	R_{max} GFlop/s	N_{max} Order	$N_{1/2}$ Order	R_{Peak} GFlop/s
MVS-5000BM Cluster IBM JS20 (dual IBM PowerPC 970 - 1.6 GHz w/Myrinet)	330	1401	280000	45000	2112
Cray X-1 (800 MHz)	120	1400.4	230400	26496	1536.0
IBM xSeries 335 cluster (dual 3.06GHz Xeon w/InfiniBand)	384	1389	120000		2350
IBM eServer pSeries 690 Turbo(1.3 GHz Power 4 w/Colony)	512	1384.0	200000		2662
NEC SX-7/160M5(1.81ns)	160	1378	200000	15200	1412.8
Dell PowerEdge 1850 (Xeon 64 3.2GHz w/Topspin InfiniBand)	256	1349	220440	110220	1638
Intel ASCI Option Red (200 MHz Pentium Pro)	9152	1338.	235000	63000	1830
Legend DeepComp 1800 (2GHz Pentium 4 w/Myrinet)	512	1297	172000		2048
Self Made Pentium4 Xeon (80-3.06 GHz, 72-2.8 GHz, 112-2.4 GHz, 256-2.2 GHz w/GigE)	520	1283	260000		2557
Intel EM64T (2 way 3.2 GHz Intel EM64T w/Myrinet D)	256	1269	241920		1638
HP Integrity rx2600 Itanium2 (1.3 GHz w/Myrinet)	304	1253	256000		1580
IBM eServer (Opteron 2.2 GHz w/Infiniband)	400	1246	200000		1760

1258




Conclusions



- Easy to accelerate the Linpack performance of workstations and clusters using Tesla and CUDA
- Increasing the performance per node reduces the cost of high performance interconnects on clusters
- Code is available from NVIDIA

Results on workstation



1 QC Core 2 and
1 Tesla C1060 faster
then 2 Xeon